

# Managing PostgreSQL on Windows

# Outline

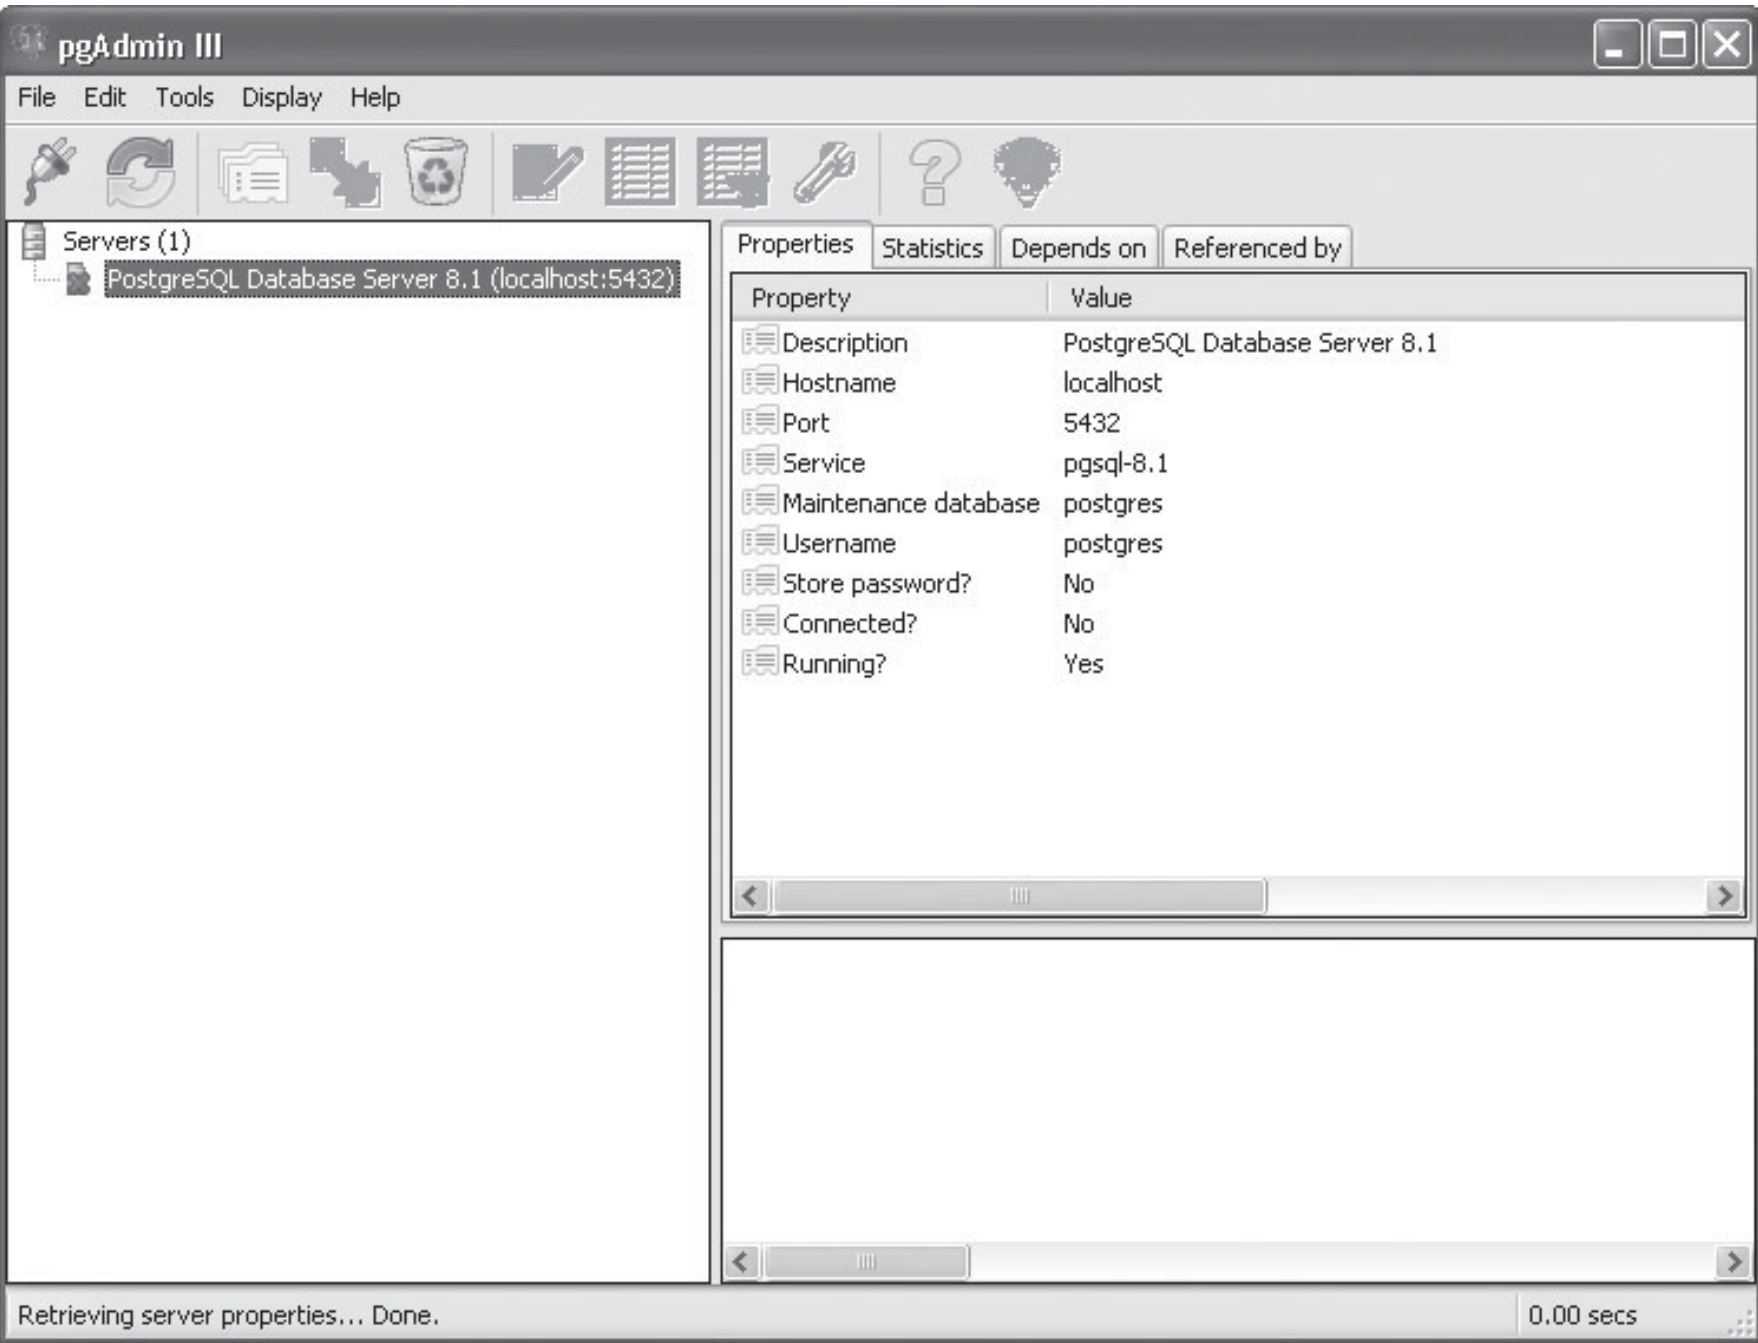
1. The pgAdmin III
2. Parts of the PostgreSQL system
3. Practices – Create a new application

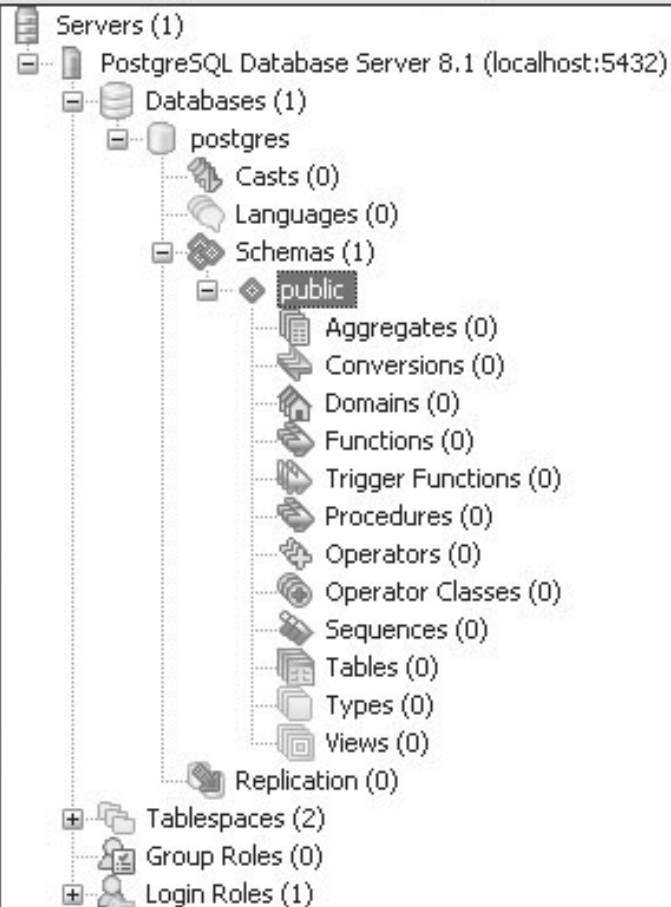
- localhost
- Port: 5432
- Account: postgres
- Password: 12345678

# 1. The pgAdmin III

# The pgAdmin III

- Any function you need to perform on your PostgreSQL system you can do from within the pgAdmin III graphical interface
- Location: `~bin\pgadmin3.exe`
- Default:
  - localhost
  - port: 5432
- Add new connect: File → Add server
- Connect server: right click → Connect





Properties

Statistics

Depends on

Referenced by

Property	Value
Name	public
OID	2200
Owner	postgres
ACL	{postgres=UC/postgres,=UC/postgres}
System schema?	No
Comment	Standard public schema

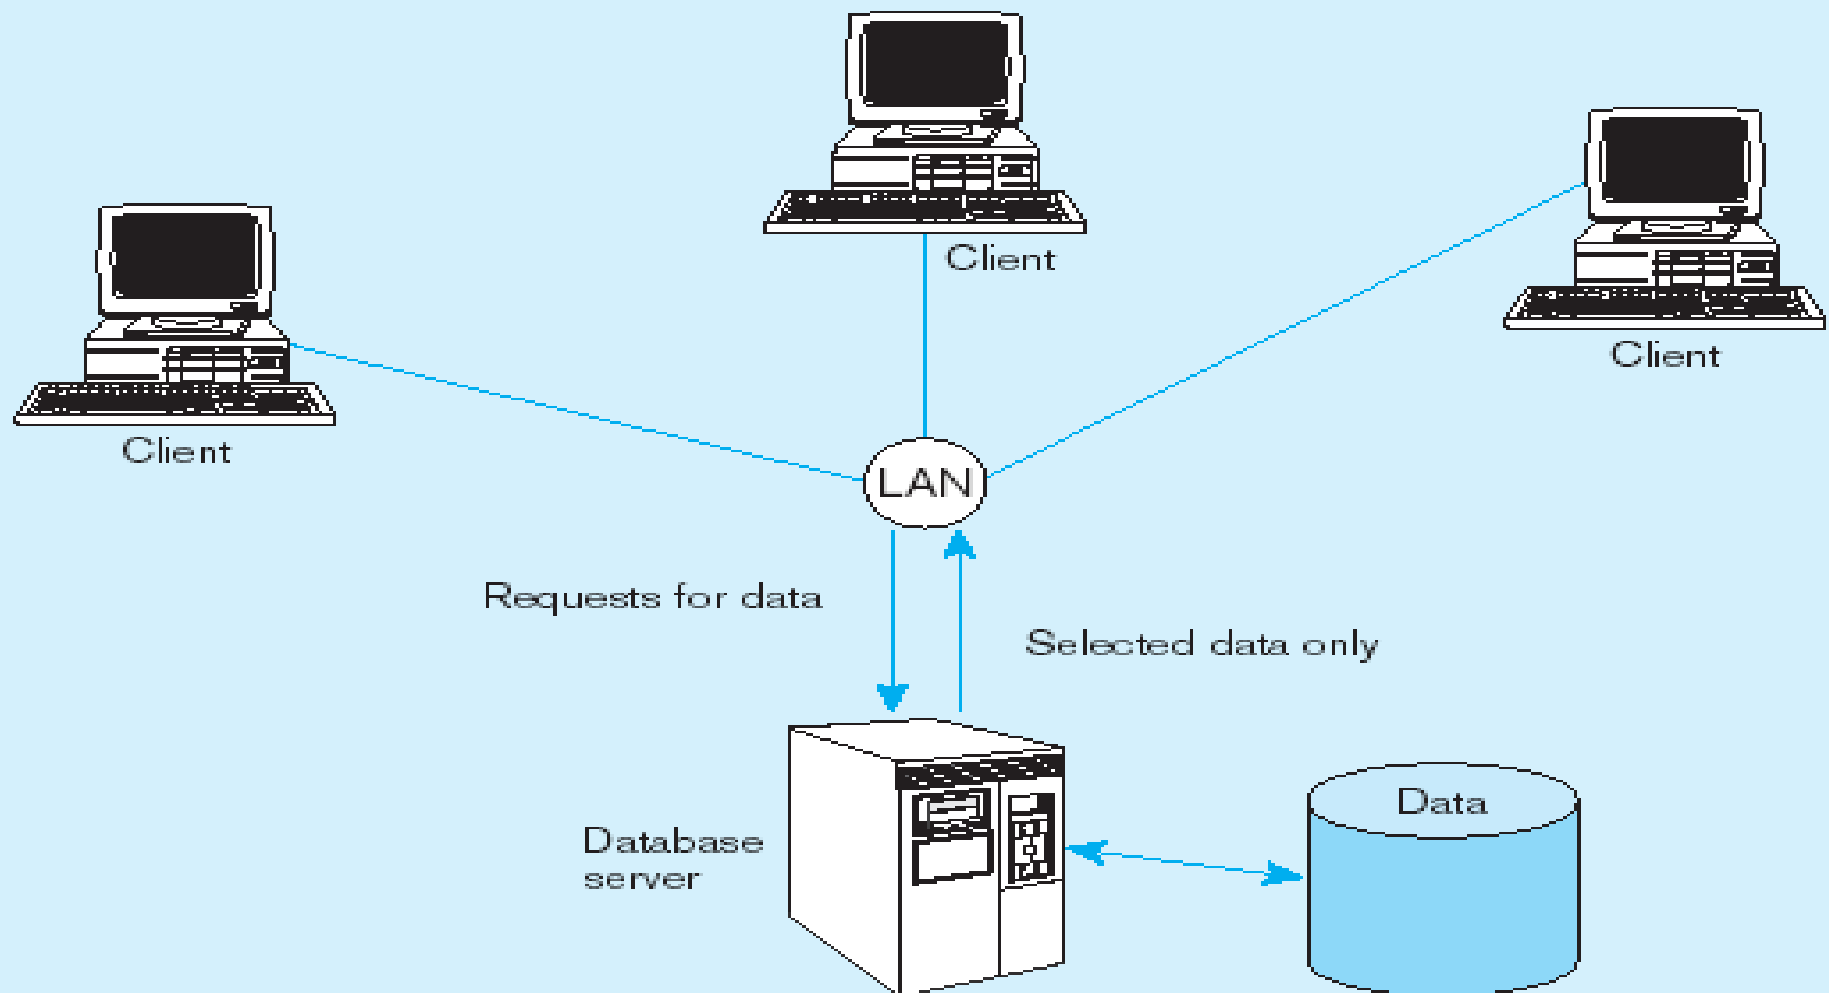
```
-- Schema: "public"
-- DROP SCHEMA public;

CREATE SCHEMA public
  AUTHORIZATION postgres;
GRANT ALL ON SCHEMA public TO postgres;
GRANT ALL ON SCHEMA public TO public;
COMMENT ON SCHEMA public IS 'Standard public schema';
```

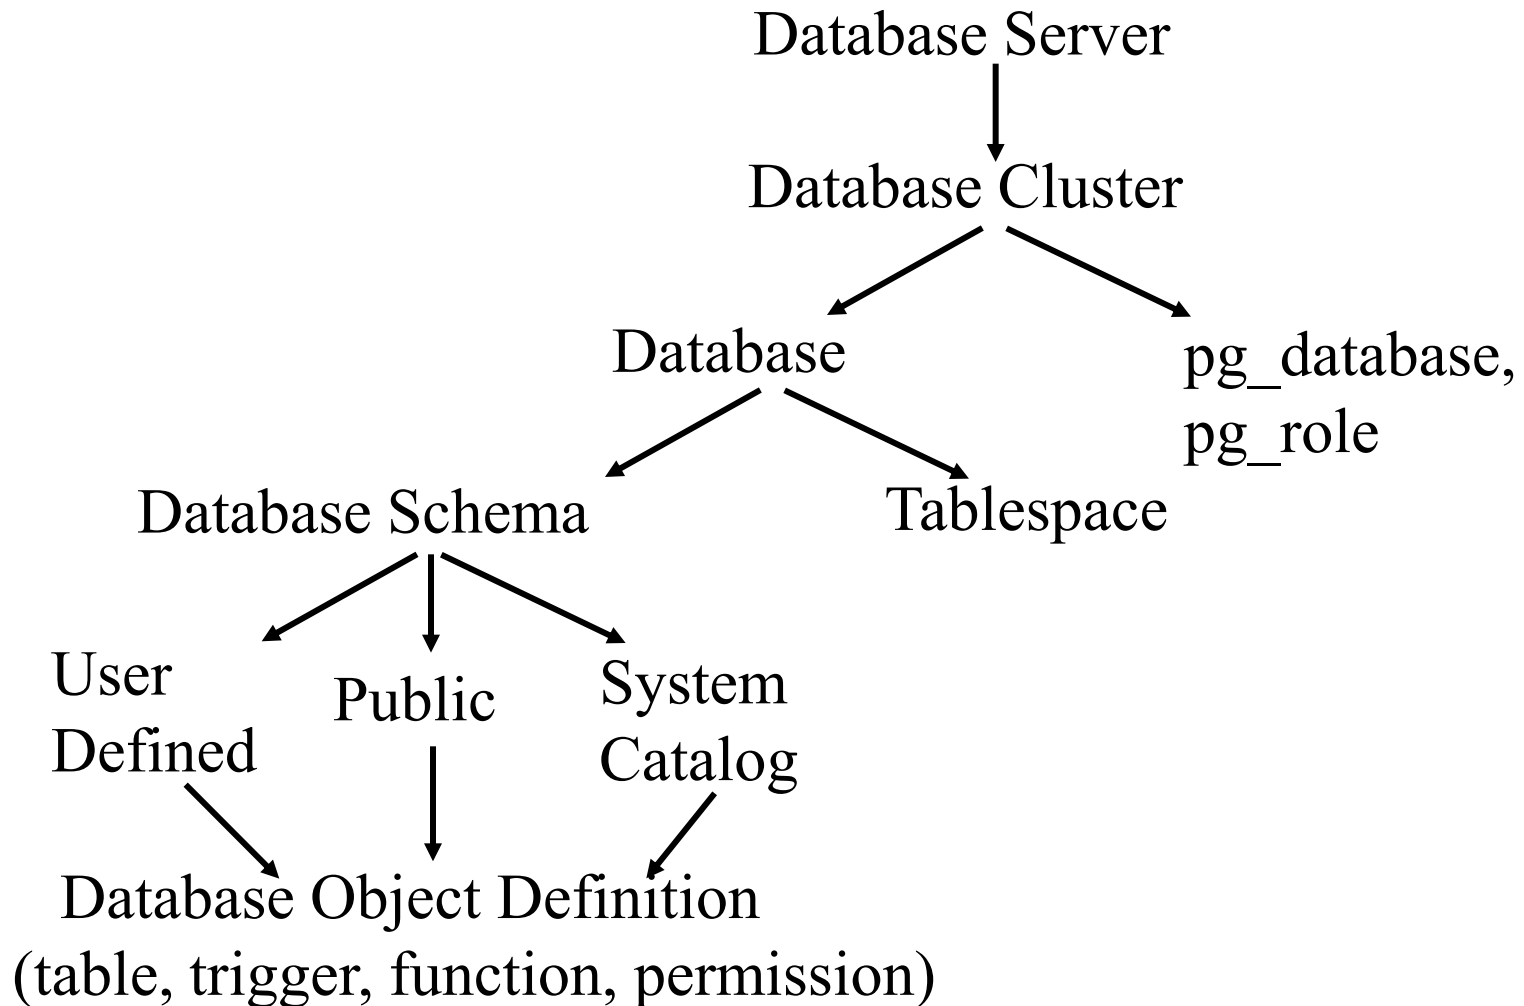
## 2. Parts of the PostgreSQL system



# Two-tier Database Server Architecture



# Database System Structure



# Database Server

- Database server
  - A computer program that provides database services to other programs (client applications) that access the server via a network
- A database server divides a client application onto:
  - A front end that runs on a user's computer and typically performs simple tasks as displaying results and
  - A back end that runs on the server computer and performs:
    - User authentication,
    - Transaction control,
    - Query optimization,
    - Database access

# Database Cluster

- A database cluster is a directory on disk where all database data will be stored (data area)
- Data in a cluster is stored as a collection of databases
- Cluster databases are managed by a single server instance
- After initialization, a database cluster contains:
  - A database named `postgres`,
  - A database named `template1`
  - A database named `template0`,
  - ...

# 5 basic components

- Tablespaces
- Databases
- Schemas (listed under each individual database)
- Group Roles
- Login Roles

# Tablespaces (1/2)

- Table spaces allow a DBA **to define locations in the file system** where the files representing database objects can be stored
- Advantages of using table spaces:
  - If the partition or volume on which the cluster was initialized runs out of space and cannot be extended, a table space can be created on a different partition and used until the system has been reconfigured
  - Table spaces allow a DBA to optimize performance by placing mission critical database objects (like indexes) on highly reliable and fast devices

# Tablespaces (2/2)

- After initialization, two default tablespaces created:
  - **pg\_default**: the default location for all database objects
  - **pg\_global**: hold PostgreSQL system catalogs, containing internal Data Dictionary information
- When new database objects are created, you must specify which tablespace are they are stored in
- Creating a new tablespace:
  - must point to an empty directory on the system → **create directory first**
  - *postgres* must have permission to write to the directory
  - default: postgres is a normal account → **grant permission to directories**

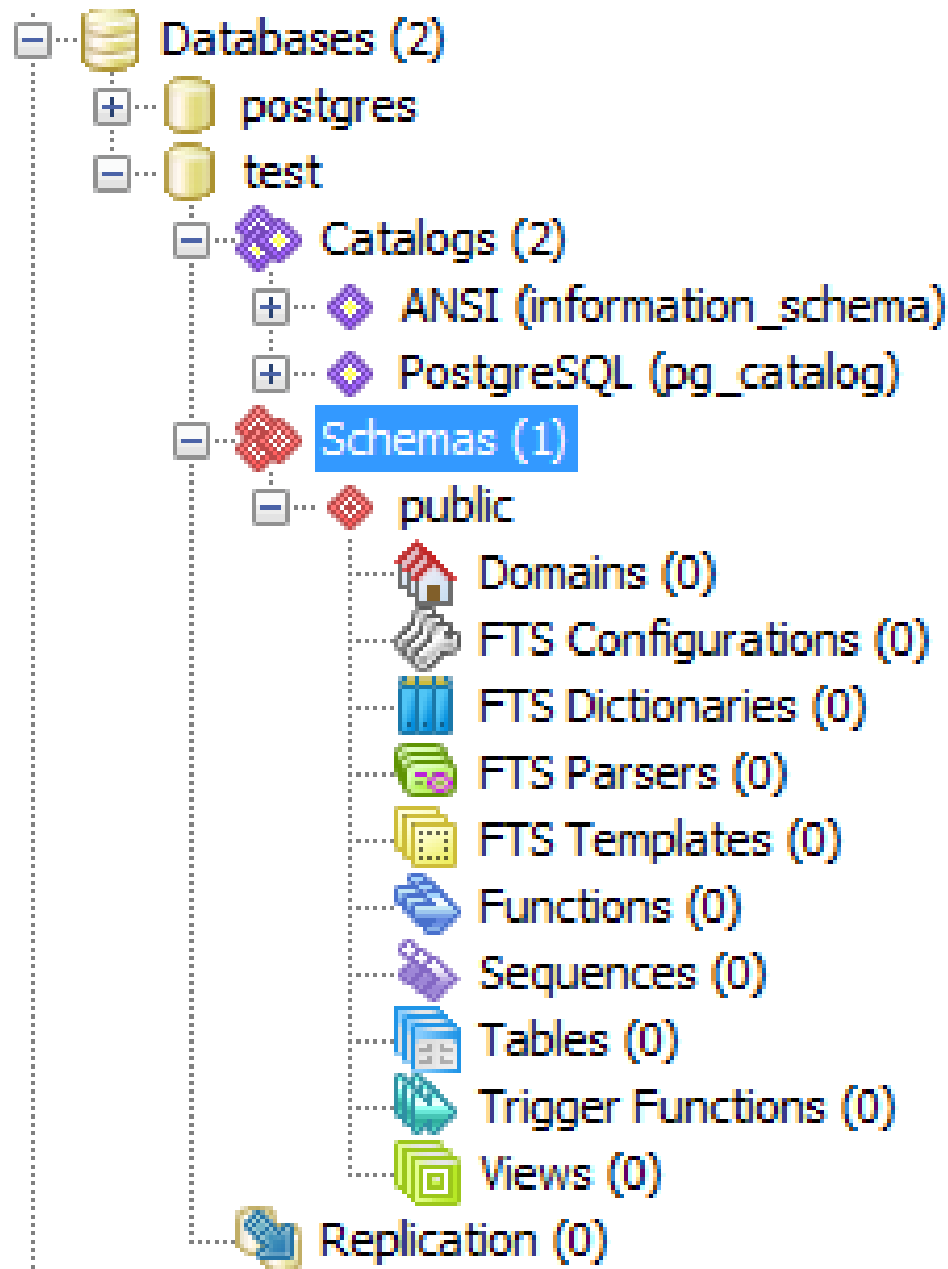
# Databases (1/2)

- The **core objects** in PostgreSQL
- Each client connection to the server can access the data in only one database
- To access data in more than one database a client must make more connections
- The default database created during the PostgreSQL installation is *postgres*:
  - contains the default system tables for handling the internal PostgreSQL Data Dictionary
- *template0* and *template1* (**NOT shown in pgAdmin III**) are used to create new databases
  - *template1* can be modified



# Databases (2/2)

- Each database object contains 4 types of objects:
  - Casts: control how Postgres casts from one datatype to another (**NOT view in pgAmin III**)
  - Languages: these are the languages you can define stored functions, aggregates and triggers in (**NOT view in pgAmin III**)
  - Schemas: **the most important objects within the database**, containing the tables, triggers, functions, views, and other objects for handling data
  - Replications: define copies (or replicas) of the PostgreSQL database in a fault-tolerant operation



# Template

- CREATE DATABASE actually works by copying an existing database
- Default, it copies the standard system database *template1*
- There is a second standard system database named *template0*
  - the same data as the initial contents of *template1*
  - *never be changed*

# Schemas (1/2)

- The **most important objects** within the database
- A database **contains one or more schemas**, which contain **database object** (table, data type, domain, function, trigger) **definitions**
- While users can only access objects **within one database at a time**, they can access **all of the schemas within that database**, *if it has permissions*
- Unlike databases, schemas are not rigidly separated

# Schemas (2/2)

Schema Object	Description
Aggregates	Defines functions that produce results based on processing input values from multiple records in a table (such as a sum or average)
Conversions	Defines conversions between character set encodings
Domains	User-defined data types
Functions	User-defined functions
Trigger Functions	User-defined table triggers
Procedures	User-defined functions that manipulate data but do not return a value
Operators	User-defined operators used to compare data
Operator Classes	Defines how a data type can be used within an index
Sequences	Defines a sequenced number generator
Tables	User-created data repositories
Types	User-defined data types used in the database
Views	User-created queries combining data from multiple tables

# Catalogs

- A pgAdmin catalog is a schema
- Hold meta data information and built-in Postgres objects
- 2 types:
  - system catalog: pg\_catalog
  - information catalog: information\_schema

# The System Catalog

- Every database system must have a **meta-database of information on the schema** which it contains.
  - The names of the relations in the schemas
  - The names of the columns of each relation.
  - The data type of each column.
  - The integrity constraints on the relations.
  - Information about indices on the relations.
  - The *access privileges* for the elements of the schema.
- This database is often called the *system catalog*.

# The System Catalog Schema

- Each PostgreSQL database contains a `pg_catalog` schema
- Normally, it is copied from the `template1` database
- The `pg_catalog` schema contains tables with information about database objects like:
  - Schemas (`pg_namespace`)
  - Tables, indexes, sequences, and views (`pg_class`),
  - Data types (`pg_type`),
  - Functions and procedures (`pg_proc`),
  - Table columns (`pg_attribute`),
  - Check, unique, primary key, and foreign key constraints (`pg_constraint`),
  - Aggregate functions (`pg_aggregate`),
  - Triggers (`pg_trigger`),
  - Planner (optimizer) statistics (`pg_statistics`), and



# Information\_schema and pg\_catalog

- The *pg\_catalog* schema is the standard PostgreSQL meta data and core schema.
- The *information\_schema* is part of the ANSI standard, but is not quite so standard. Oracle and DB2 evidently still don't support
- A lot of this ***information overlaps*** with information found in the *information\_schema* and *pg\_catalog*, but **the *information\_schema* is much easier to query**
- Although not named explicitly in the **search\_path**, *pg\_catalog* is the first schema to be searched
  - `SELECT * FROM pg_tables`
  - `SELECT * FROM pg_catalog.pg_tables`

# Group Roles

- Create access **permissions for groups of users**
- While you can grant an individual user account access directly to a database object, the preferred method is to use Group Roles
- pgAdmin III only allows you to grant Group Roles access to database objects
- Default, **public** group role:
  - applies to all users on the PostgreSQL system
  - NOT able to remove any user account from the public Group Role
  - does not appear in the pgAdmin III Group Roles listing

# Login Roles (or user accounts)

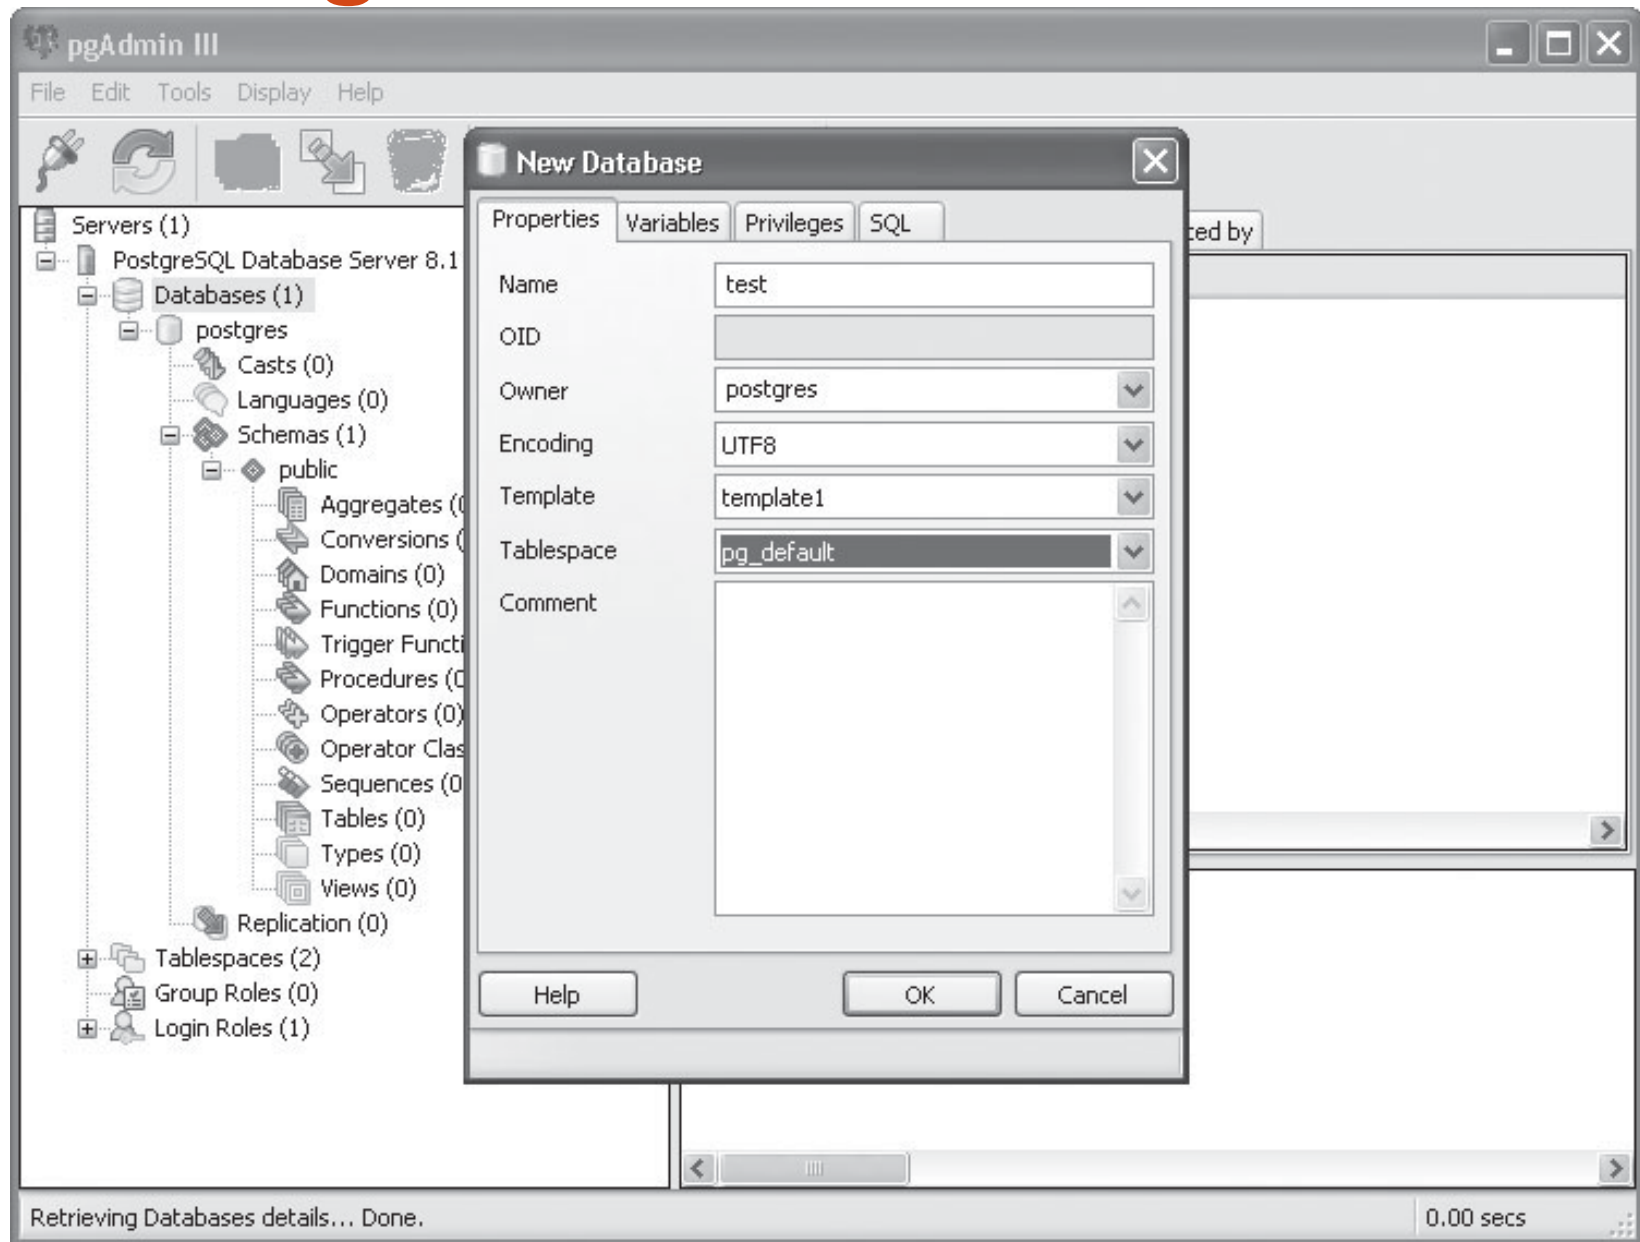
- Are roles that are allowed to log into the PostgreSQL server
- **Each database user** should have an **individual account** for logging into the PostgreSQL system
- That account is then **assigned as a member** of the appropriate **Group Roles that grant privileges** to the database objects required
- Allows you to **easily change access for database objects** without having to touch hundreds (or even thousands) of individual user Login Roles

### 3. Practices – Create a new application

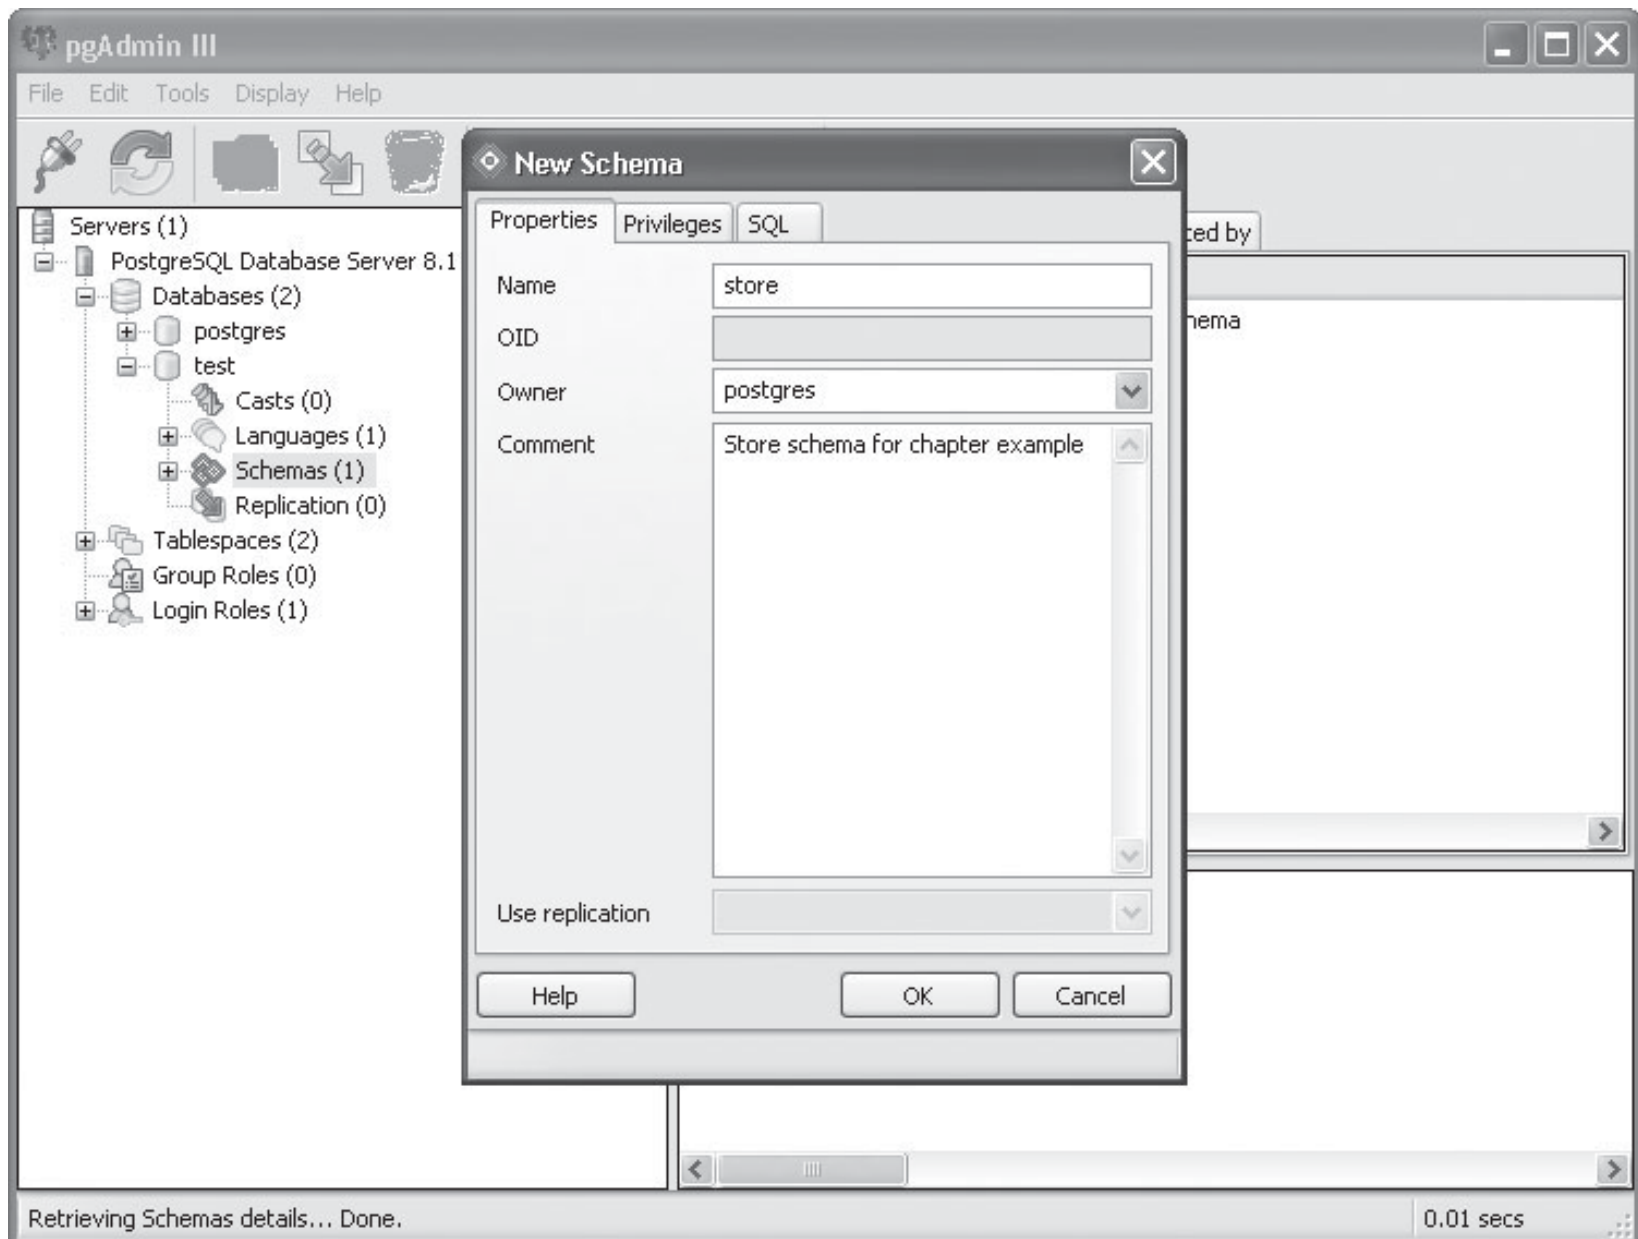
# Practices – Create a new application

- Create a database ***test***
  - Customer
  - Product
  - Order
- Create two Group Roles
  - *Salesman* Group Role: write permission on the Customer and Order, only read permission on the Product
  - *Accountant* Group Role: write permission on the Product and Order, read permission on the Customer
- Create two Login Roles
  - salesman - Barney
  - accountant - Fred

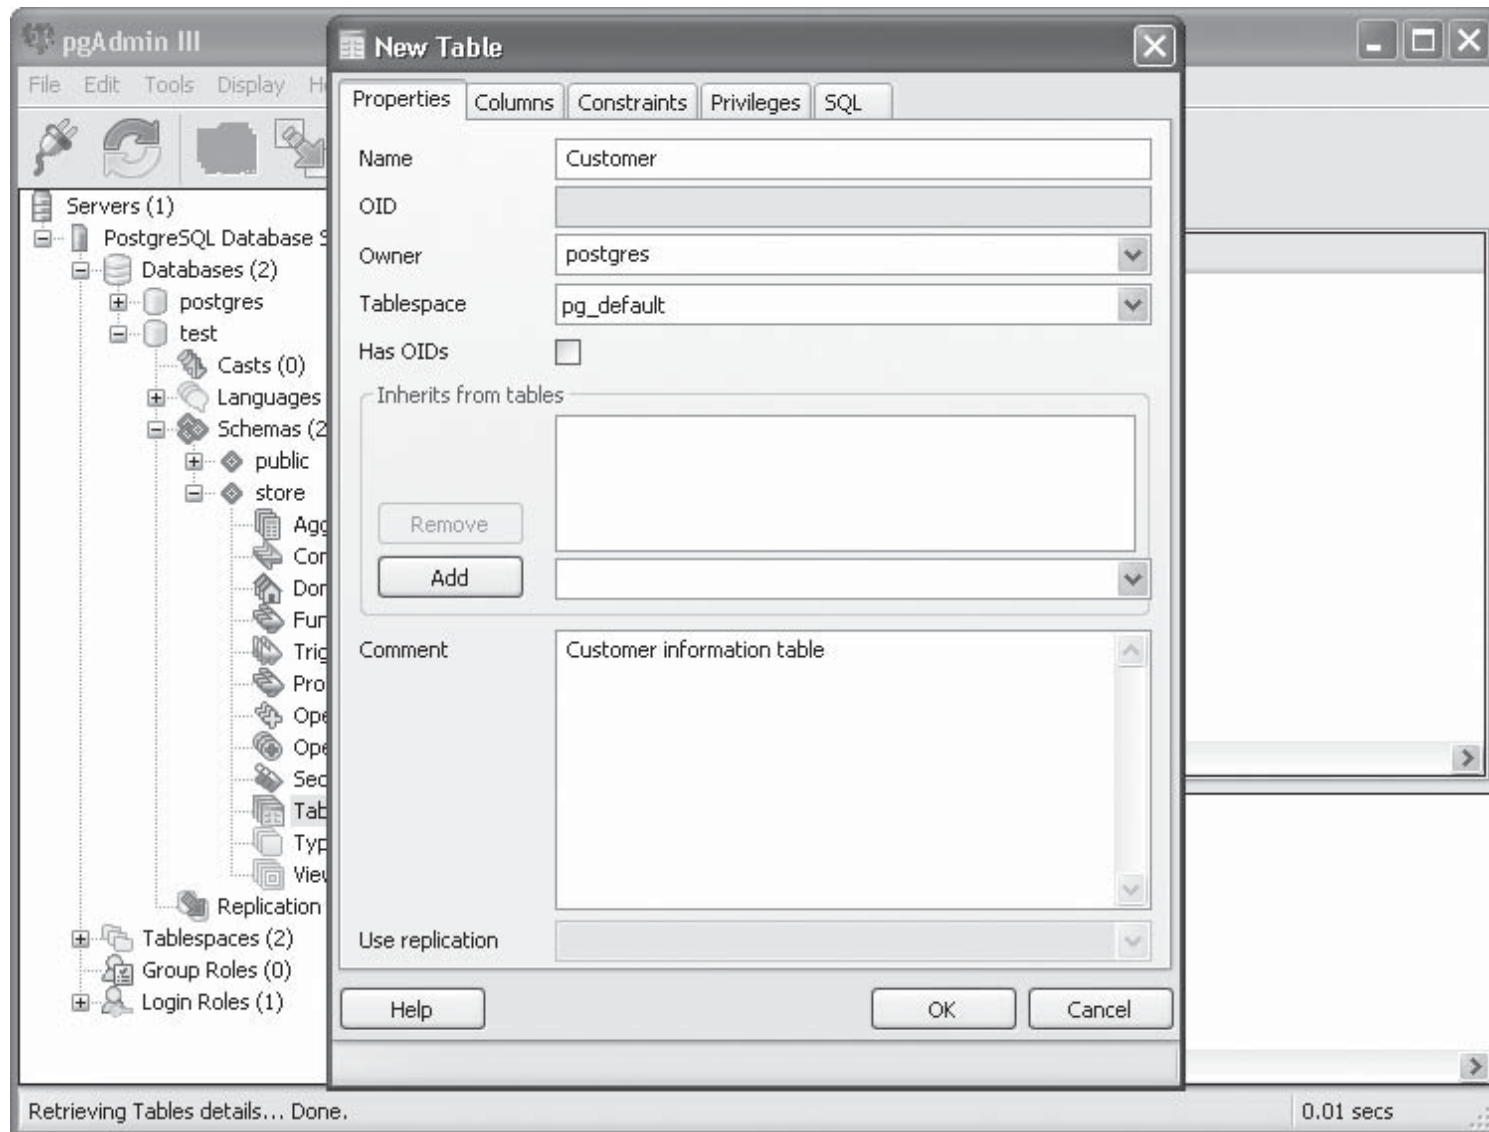
# Creating a New Database



# Creating a New Schema



# Creating the Tables





# Customer Table Columns

Column	Data Type	Description
CustomerID	char—six characters	Unique identifier for each customer
LastName	varchar	Last name of customer
FirstName	varchar	First name of customer
Address	varchar	Street address of customer
City	varchar	City of customer
State	char—two characters	State of customer
Zip	char—five characters	Postal ZIP code of customer
Phone	varchar	Phone number of customer



# Common PostgreSQL Data Types

Name	Aliases	Description
<code>bigint</code>	<code>int8</code>	signed eight-byte integer
<code>bigserial</code>	<code>serial8</code>	autoincrementing eight-byte integer
<code>bit [ (n) ]</code>		fixed-length bit string
<code>bit varying [ (n) ]</code>	<code>varbit</code>	variable-length bit string
<code>boolean</code>	<code>bool</code>	logical Boolean (true/false)
<code>box</code>		rectangular box on a plane
<code>bytea</code>		binary data ("byte array")
<code>character [ (n) ]</code>	<code>char [ (n) ]</code>	fixed-length character string
<code>character varying [ (n) ]</code>	<code>varchar [ (n) ]</code>	variable-length character string
<code>cidr</code>		IPv4 or IPv6 network address
<code>circle</code>		circle on a plane
<code>date</code>		calendar date (year, month, day)
<code>double precision</code>	<code>float8</code>	double precision floating-point number (8 bytes)
<code>inet</code>		IPv4 or IPv6 host address
<code>integer</code>	<code>int</code> , <code>int4</code>	signed four-byte integer
<code>interval [ fields ] [ (p) ]</code>		time span
<code>json</code>		JSON data

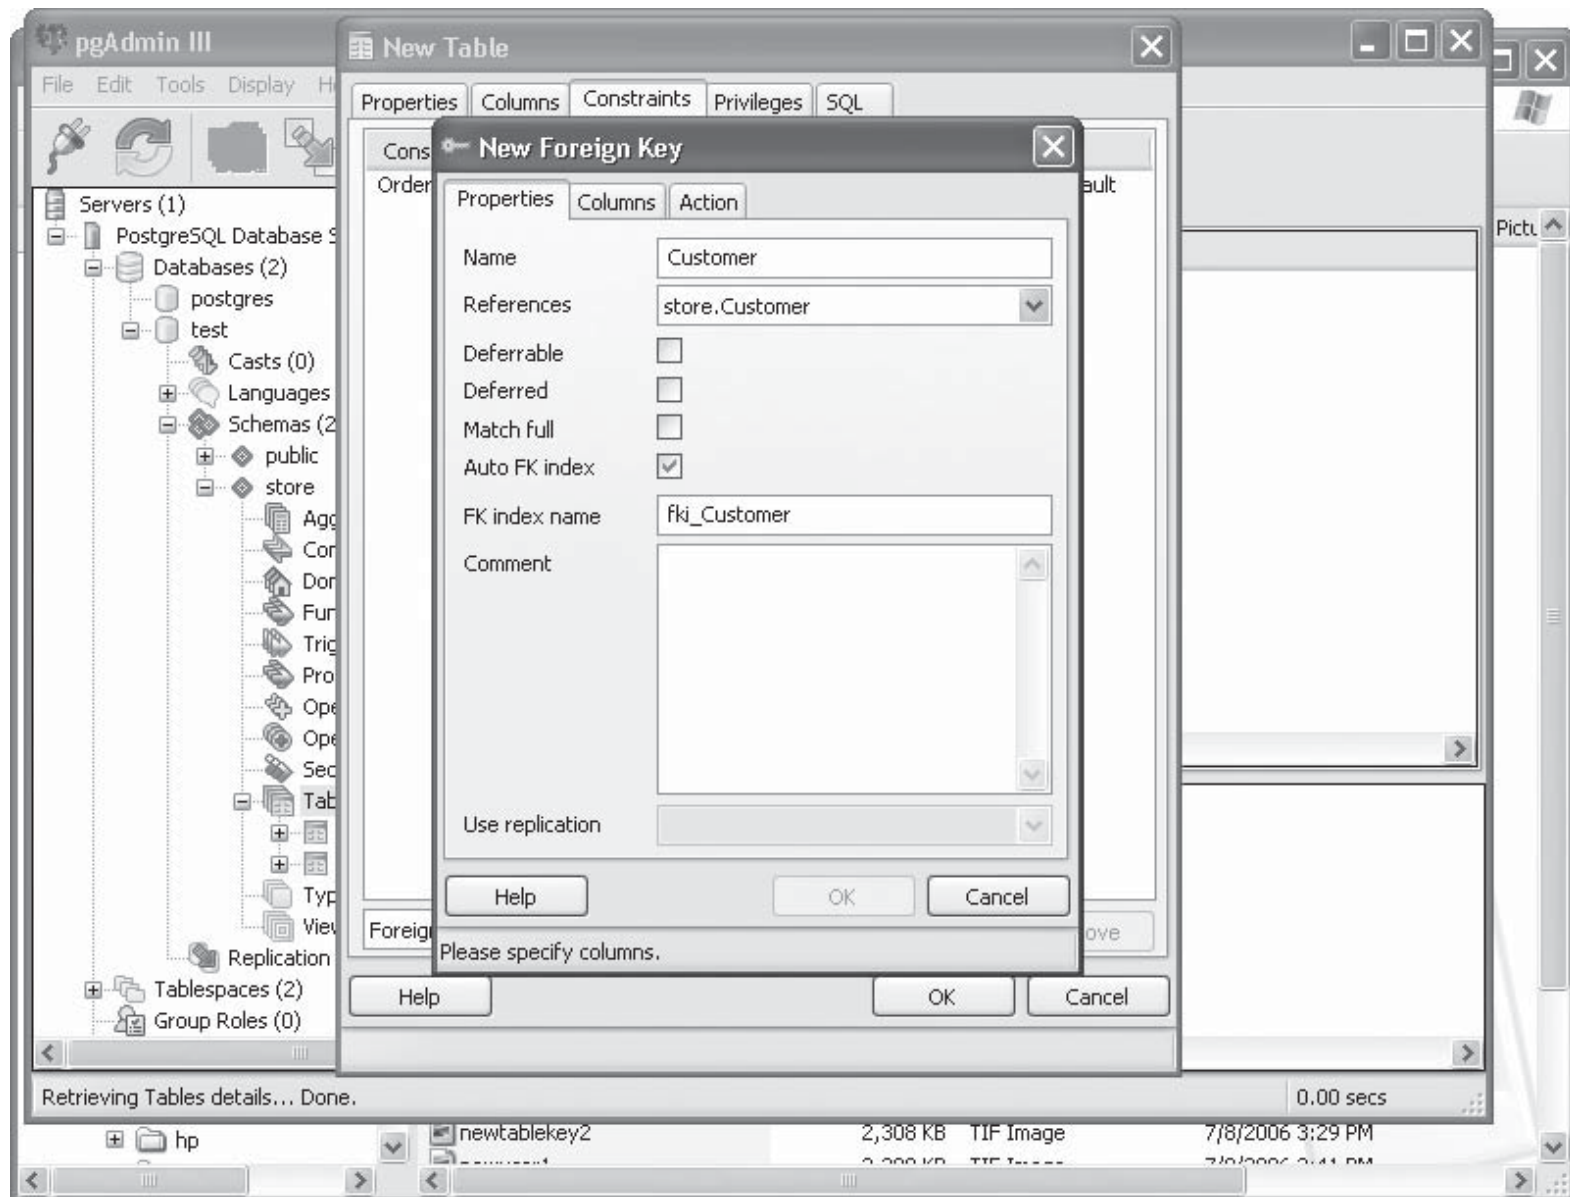
# The Product Table Columns

Column Name	Data Type	Description
ProductID	char—six characters	Unique primary key identifier that is not NULL
ProductName	varchar	Name of the product
Model	varchar	Product model number
Manufacturer	varchar	Name of the manufacturer
UnitPrice	money	Current price of product
Inventory	int4	Number of units in inventory

# The Columns for the Order Table

Column Name	Data Type	Description
OrderID	char—six characters	Unique primary key identifier that is not NULL
CustomerID	char—six characters	The CustomerID from the Customer table (not NULL)
ProductID	char—six characters	The ProductID from the Product table (not NULL)
PurchaseDate	date	Date of purchase
Quantity	int4	The number of items purchased
TotalCost	money	The total cost of the purchase

# New Foreign Key window for the Order table



# Entering and Viewing Data

pgAdmin III

File Edit Tools Display Help

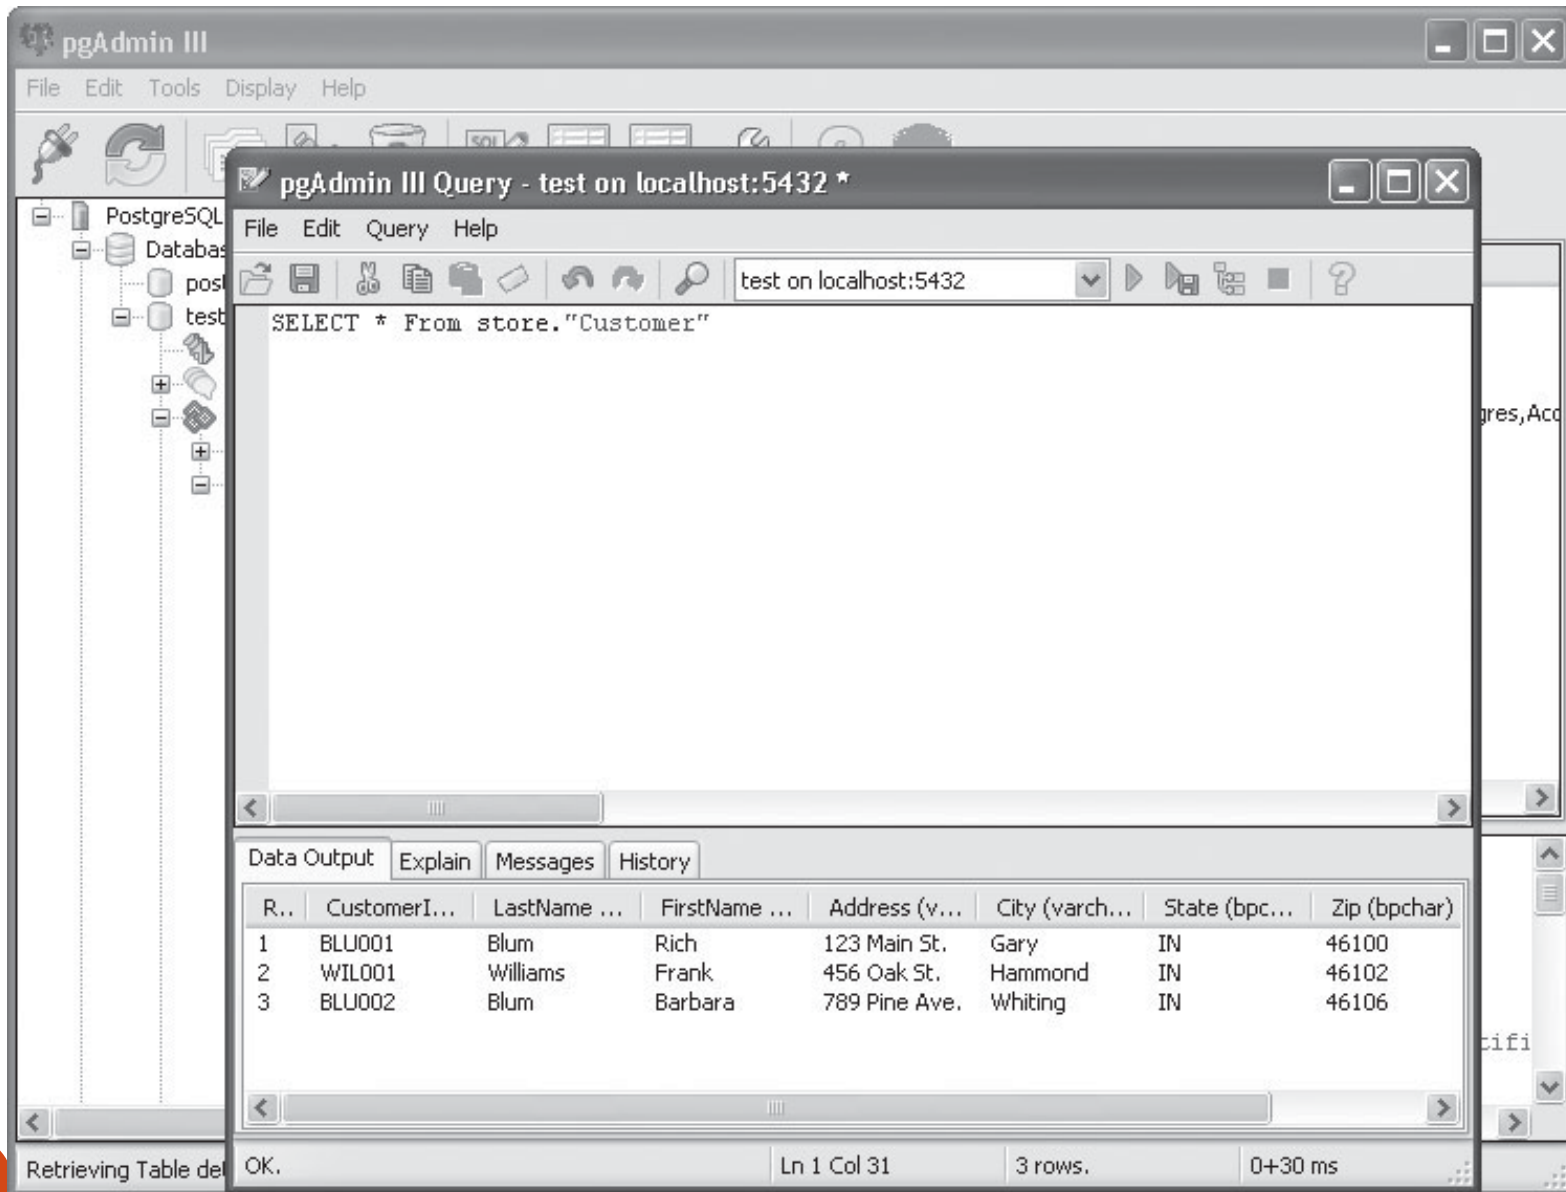
pgAdmin III Edit Data - PostgreSQL Database Server 8.1 (localhost:5432) - test - store.Customer

	CustomerID [PK] bpchar	LastName varchar	FirstName varchar	Address varchar	City varchar	State bpchar	Zip bpchar	Phone varchar
1	BLU001	Blum	Rich	123 Main St.	Gary	IN	46100	555-1234
2	WIL001	Williams	Frank	456 Oak St.	Hammond	IN	46102	555-9876
*								

Retrieving 2 rows.

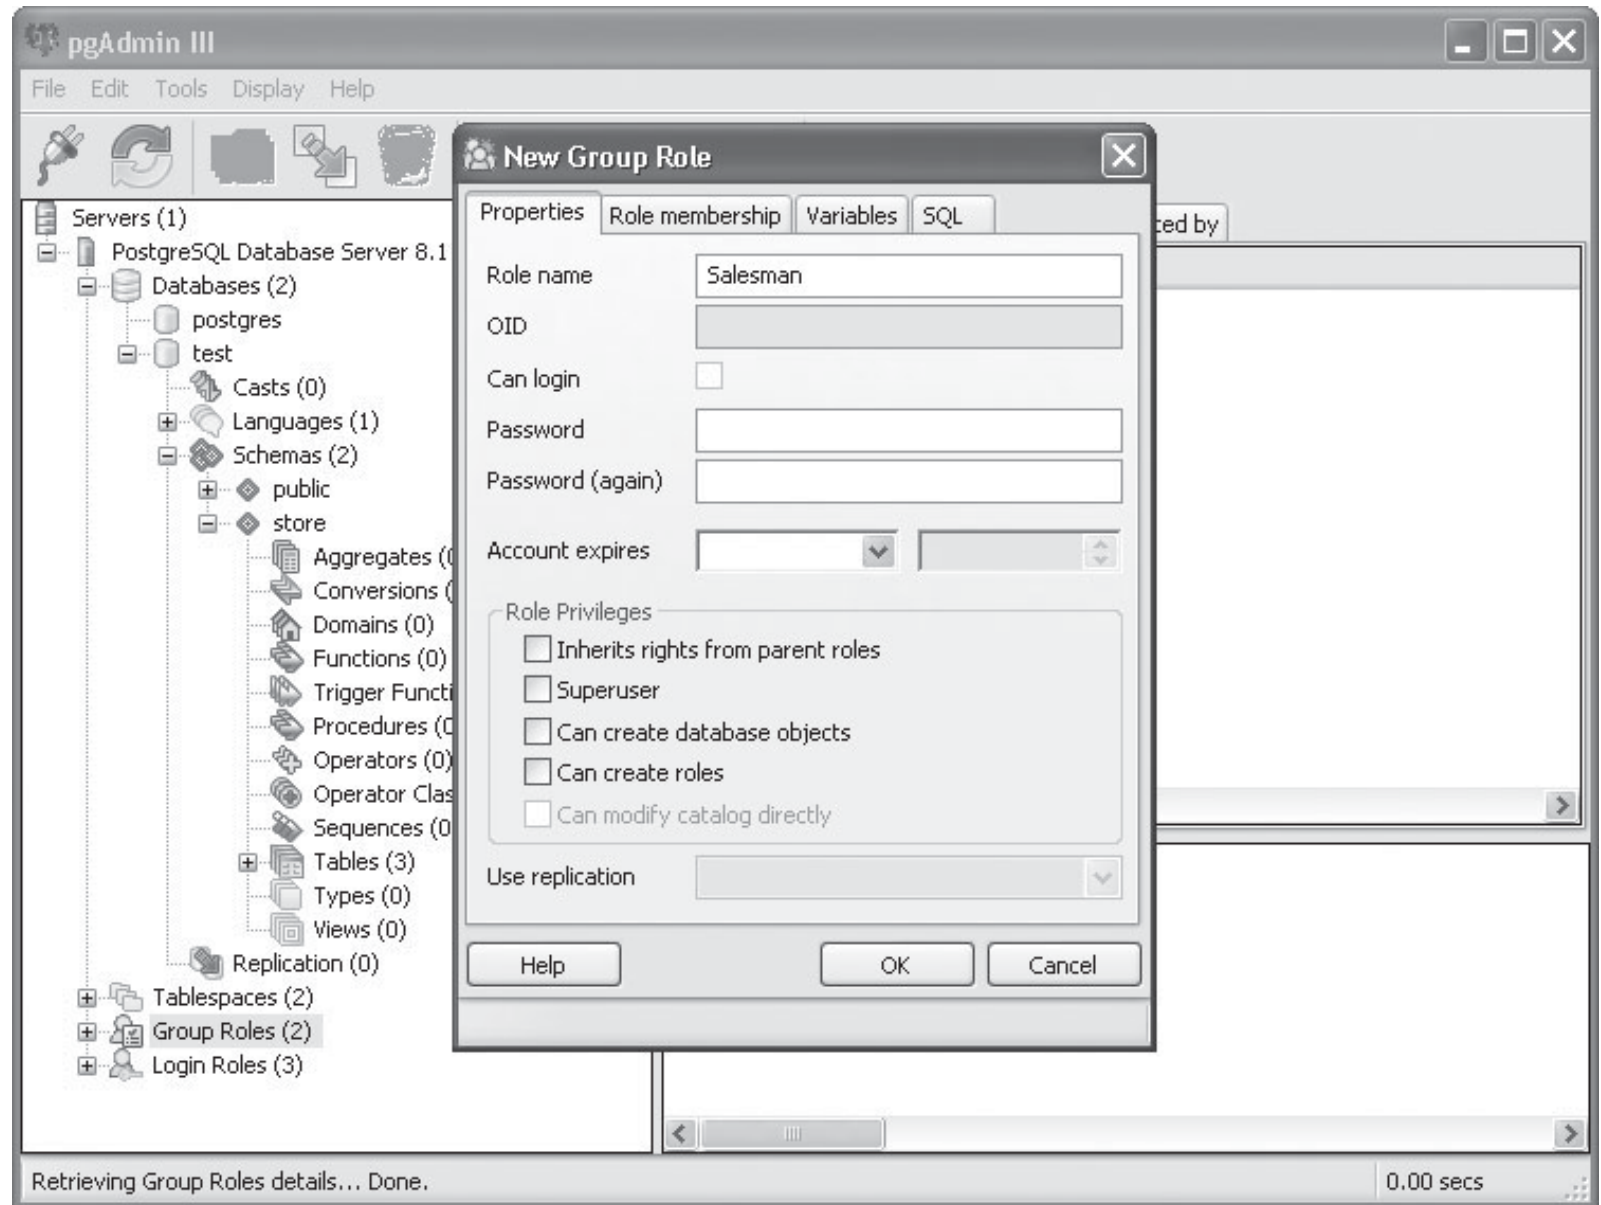


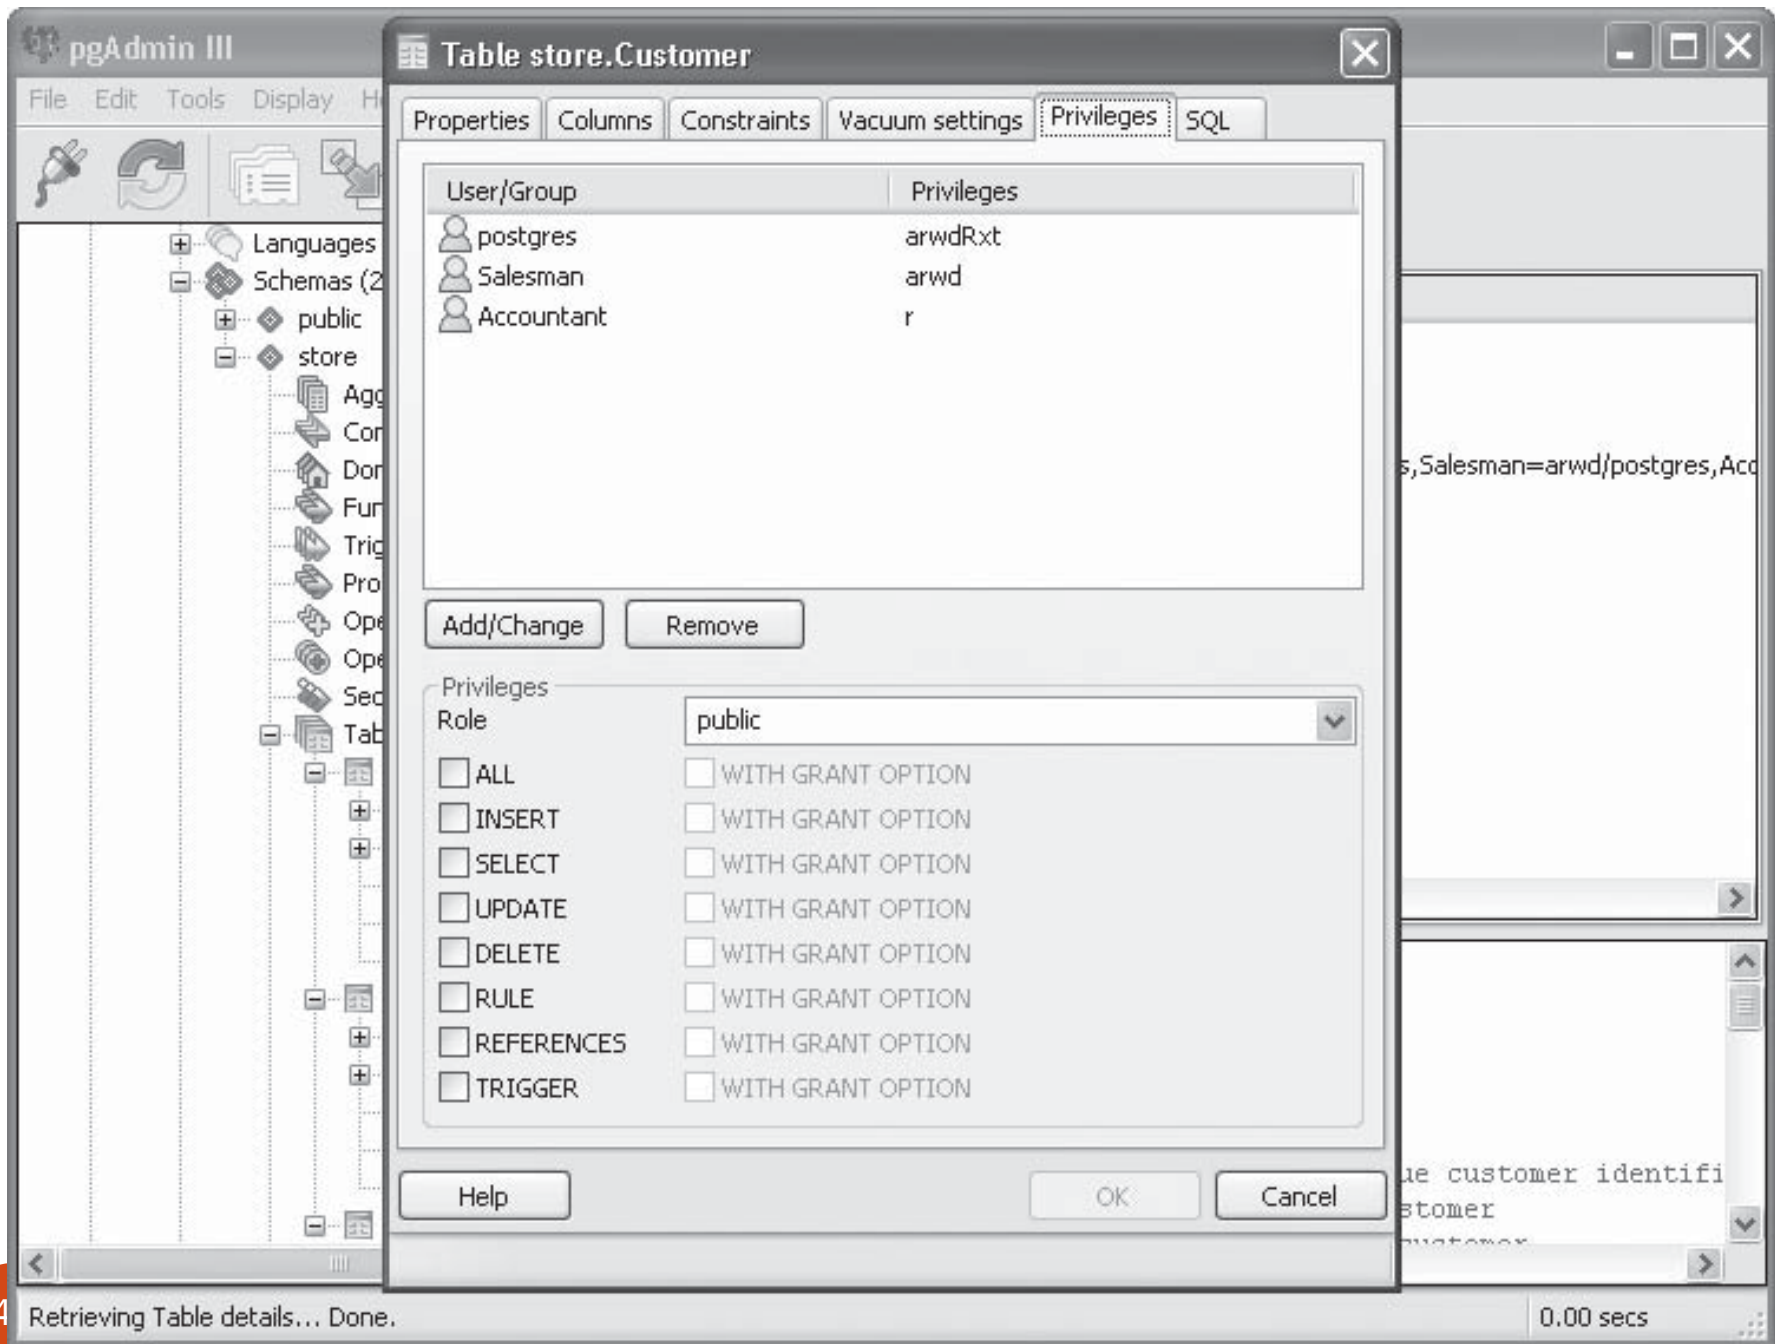
# THE pgADMIN III QUERY TOOL





# WORKING WITH USER ACCOUNTS

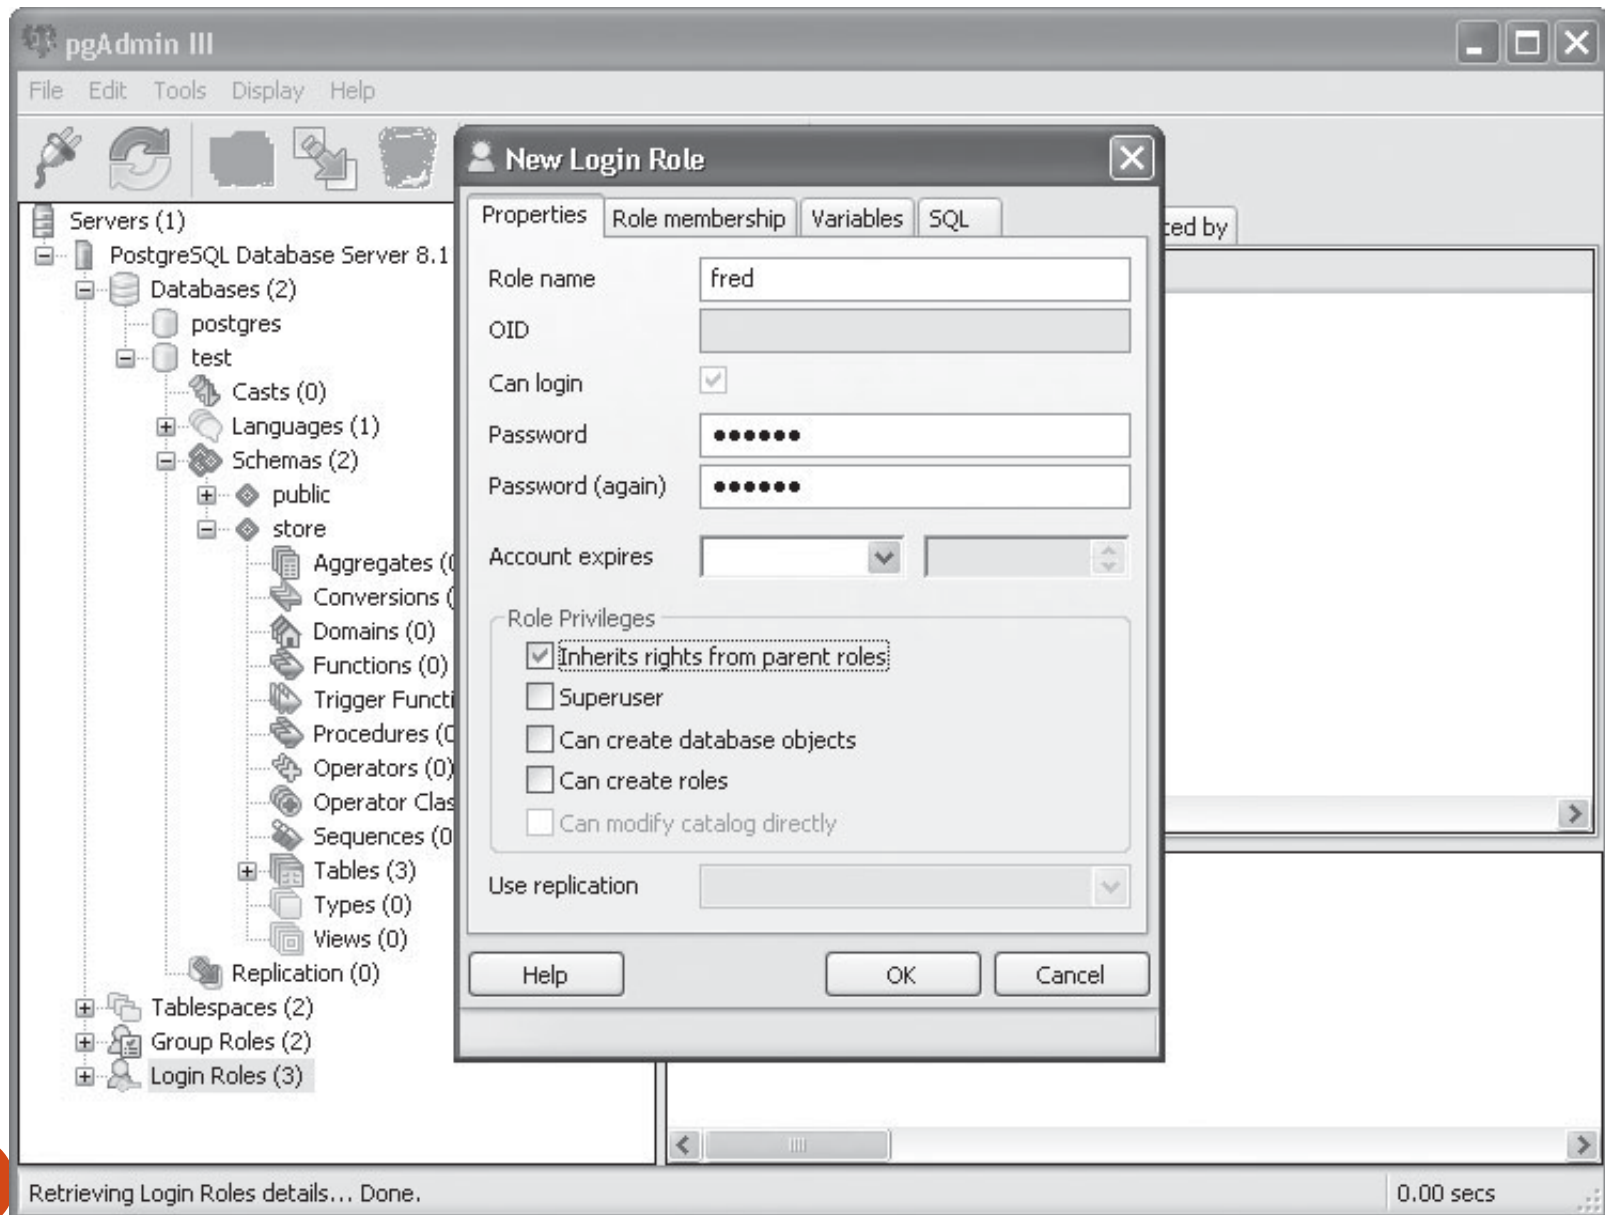




# pgAdmin Object Privilege Codes

Code	Privilege
a	INSERT (append)
r	SELECT (read)
w	UPDATE (write)
d	DELETE
R	RULE
x	REFERENCES
t	TRIGGER
X	EXECUTE
U	USAGE
C	CREATE
T	TEMPORARY

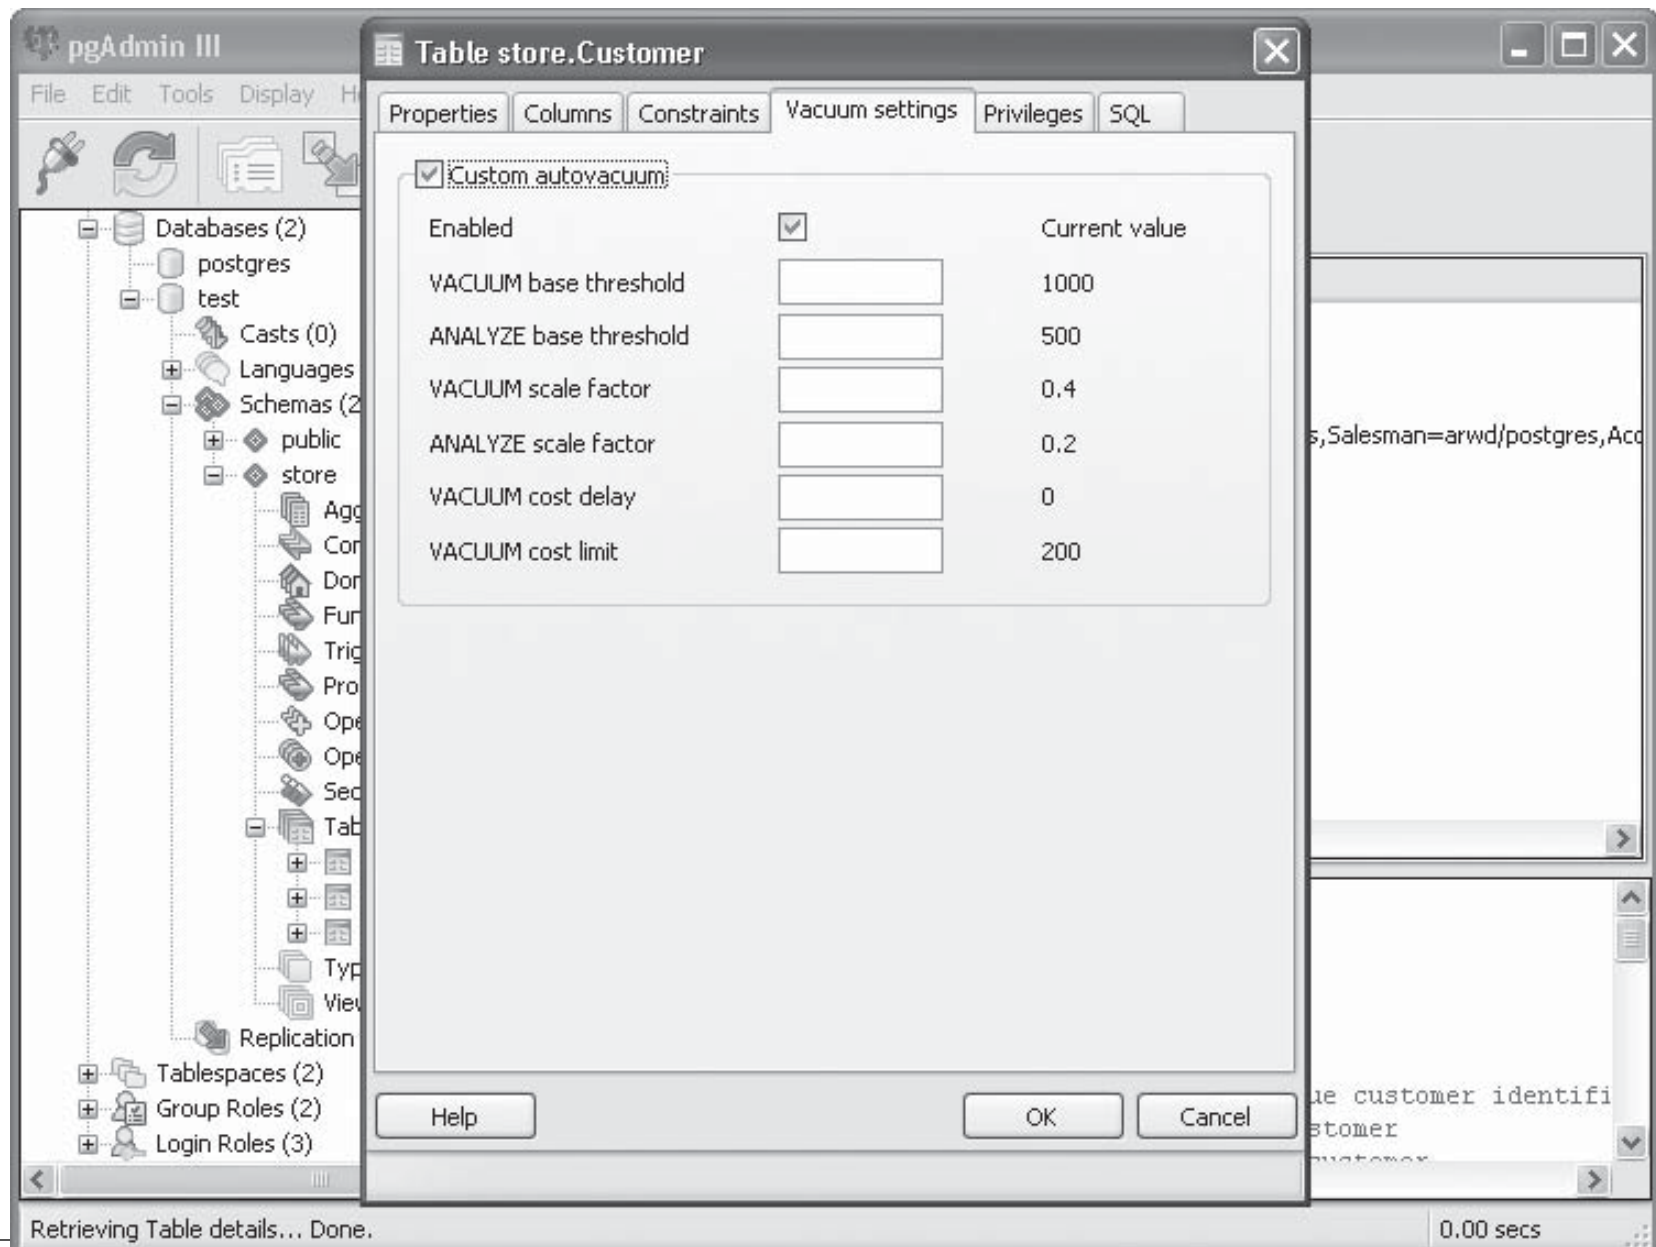
# Creating Login Roles



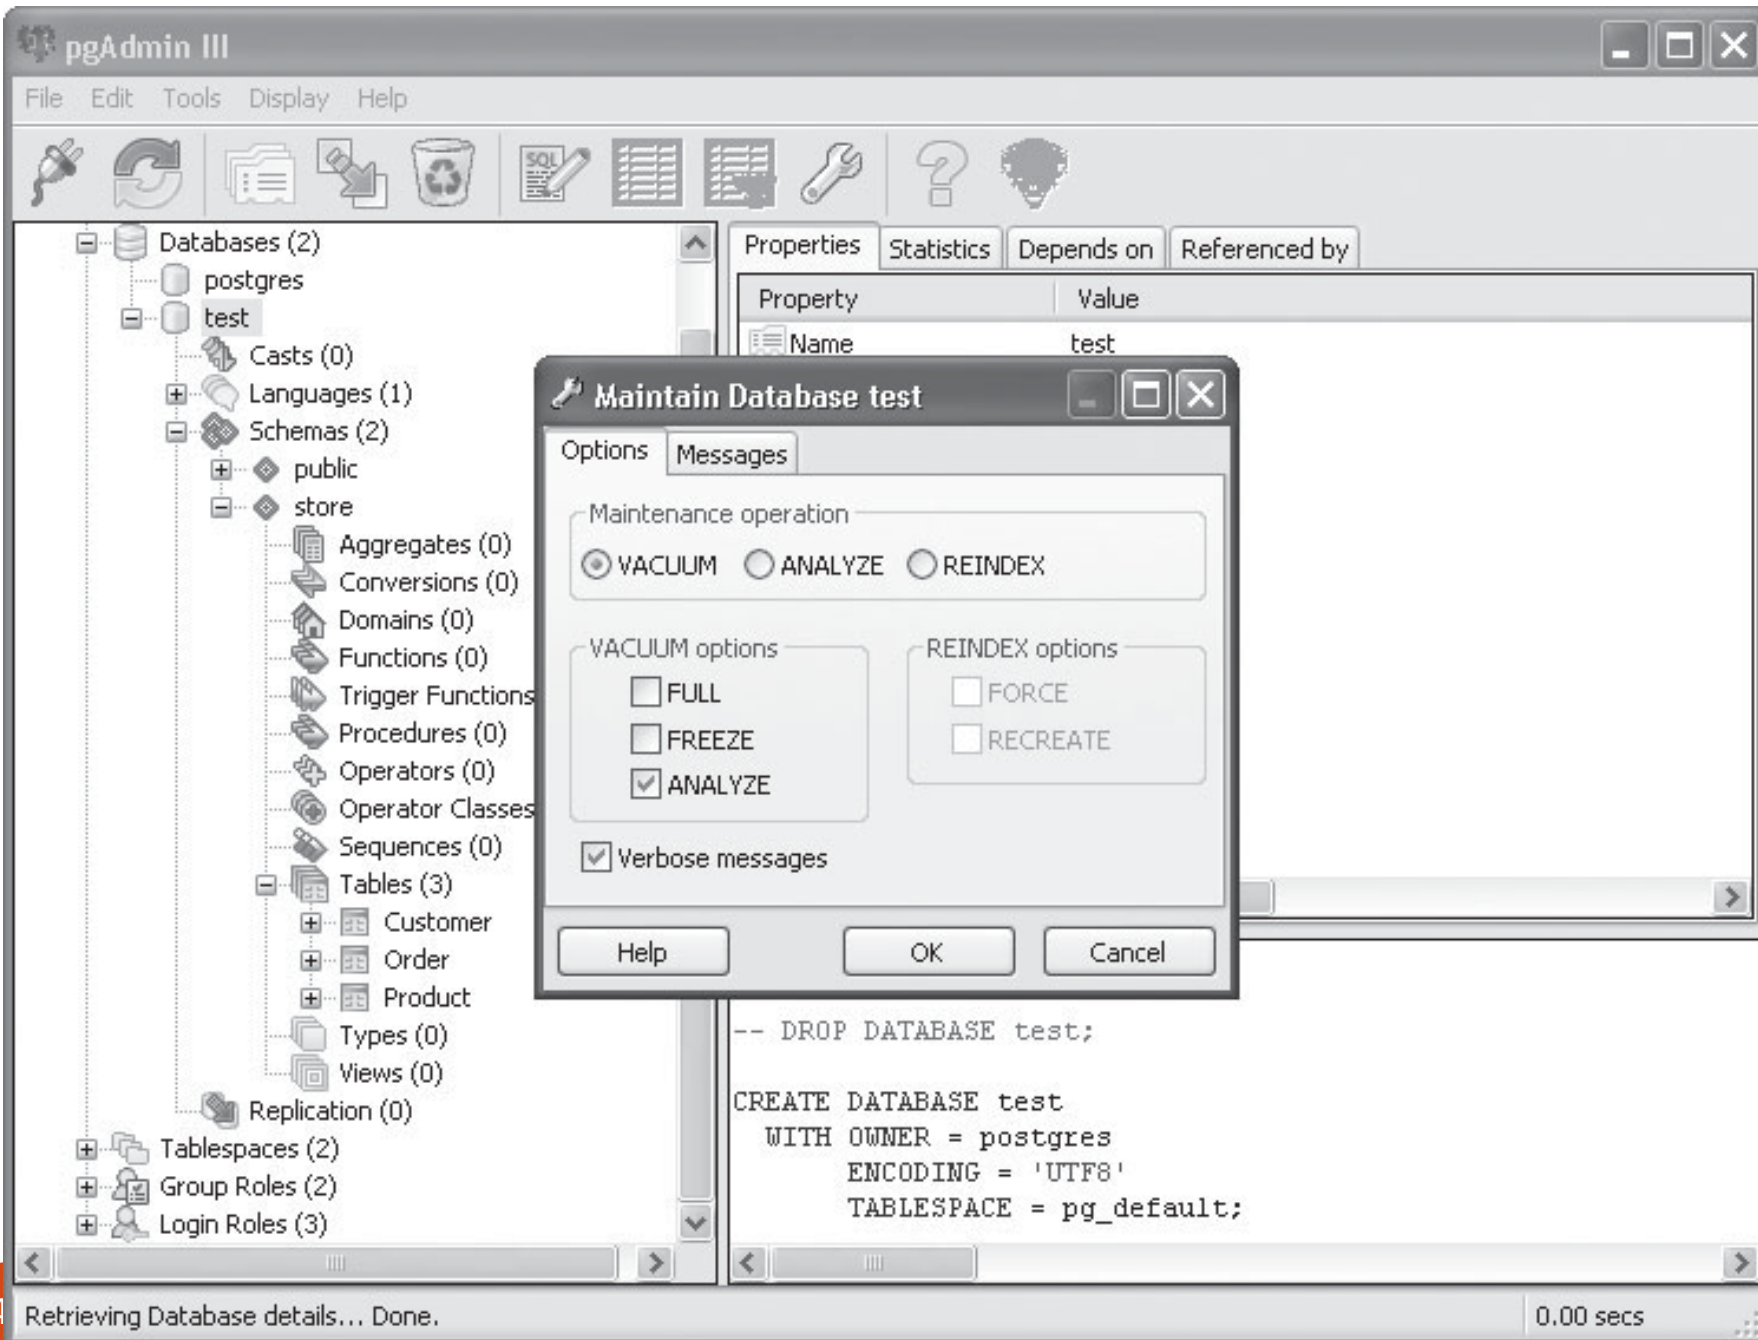
# Testing

- log in using the test database and the fred Login Role
  - psql test fred
- test=>INSERT into store."Product" VALUES ('LAP001', 'Laptop', 'TakeAlong', 'Acme', '500.00', 100);
- test=>
- test =>INSERT into store.  
"Customer"("CustomerID", "LastName", "FirstName")VALUES ('Cus001', 'Thi Oanh', 'Nguyen');

# DATABASE MAINTENANCE







# BACKUPS AND RESTORES

